Harsh Devisha

Los Angeles, CA | Permanent Resident | Open to Relocation

+1-310-696-9659 | harshdevisha91@gmail.com | harshdevisha.com | github.com/lostcache | linkedin.com/in/harsh-devisha

Backend | Systems | Developer Tools

Education

Master of Science in Computer Science - 3.9/4.0 2024 – 2026
Loyola Marymount University

Bachelor of Technology in Electrical Engineering 2017 – 2021
Pandit Deendayal Petroleum University

Experience

Founding Engineer May 2025 – December 2025
SetFindr (Early-Stage P2P Marketplace)

- Built a full-stack P2P marketplace using Next.js, Node.js, Python FastAPI, MongoDB, Qdrant, GCP, AWS and Docker.
- Implemented unstructured document indexing pipeline using GPT-4o and text-embedding-3-large models for semantic search and personalized recommendations, improving search relevance.
- Implemented image intelligence using GPT-4o to extract location details from images, reducing listing creation time by 10x.
- Designed and optimized vector search using Qdrant with HNSW indexing, achieving sub-20ms p95 latency across 10K+ embeddings.

Data Specialist June 2024 – May 2025
LMU Career and Professional Development Department

- Built data processing pipelines and analytics dashboards to track student engagement.
- Conducted cohort analysis focusing on underrepresented demographics, improving targeted outreach conversion rates by 200%.

Software Engineer - I December 2021 – December 2023
Nexuslink Services

- Built React-Redux SPA for an Ecommerce website that fired 200+ API calls on initial page load; memoized expensive computations with useMemo to skip redundant recomputes, improving Time to Interactive (TTI) by 15%.
- Created standardized API client abstraction with centralized error handling, reducing duplicated calls across 150+ endpoints.
- Automated complex document generation using headless browser (Puppeteer) for an internal tool, saving 100+ operational hours monthly.
- Implemented JWT authentication with session management, refresh token rotation, and SMTP account recovery workflows.

Open-Source Contribution

TiDB (PingCAP) | *Go* github.com/pingcap/tidb

- Cloud-Native, Distributed SQL Database (40k+ stars) with ACID, Raft, and MySQL compatibility
- Resolved #60940 (9 merged PRs): stripped volatile, sequentially-allocated plan-IDs making query-plan result files immune to manual updates for optimizer's plan-allocation order changes, cutting test-maintenance churn and eliminating flaky CI failures.

Projects

Dunno — Knowledge Graph for AI Coding Assistants | *Rust, SurrealDB, Axum, Svelte* github.com/lostcache/dunno

- Lets teams of humans and AI coding agents share plans, tasks, and file-level codebase knowledge through a graph database, eliminating the need to commit Markdown context files into source repos.
- Engineered a Rust service exposing a Clap-based CLI (JSON output for scripting) and an Axum HTTP server over SurrealDB, with layered configuration supporting embedded, local, and cloud backends for both single-user and concurrent multi-user workflows.
- Modeled developer knowledge (plans, tasks, skills, personas, mistakes, style guides, security notes) as graph entities for deterministic, structure-aware context retrieval, and built a Svelte/TypeScript interactive graph-visualization UI.

MagicTalesAI — Google Gemini, UCLA Hackathon | *Python, FastAPI, React, GCP* github.com/lostcache/MagicTalesAI

- Generates Immersive and Interactive Audiobooks from PDF/EPUB using Google Gemini Models.
- Uses Gemini 2.5 Flash for scene segmentation and TTS character dialog generation.
- Uses Google Lyria for emotion-aware background music and ElevenLabs to clone custom voices.

Technical Skills

Languages: Go, Python, TypeScript, JavaScript, Rust, C++
Backend/Frontend: Node.js, Flask, FastAPI, RESTful APIs, React.js, Next.js
Databases/Cloud: PostgreSQL, MongoDB, Redis, Qdrant, AWS (EC2, S3), Google Cloud
Tools: Git, Jira, Click-Up